



# Next generation of Tencent OLAP Engine

Tencent TEG LongYue



数据平台部

# CONTENTS

- **01** | Background
- **02** | Storage: Various Columns and Indexes
- **03** | Computation: Integrated with Presto
- **04** | Benchmark and Applications

# Background

*MercsDB*

- **Data :**

1. Thousand Columns, 10 Billions Rows
2. Index on arbitrary column(s)
3. Real-Time Write
4. Both Row-oriented and Column-oriented
5. Different Indexes

*ElasticSearch*

! =


*Presto/Impala*

- **Performance :**

1. Second response on 10 Billions rows query
2. MPP: both ad-hoc query and real-time query
3. Real-time Write: 100 Billion rows / day

*ClickHouse*

# History



# Current Status

Clusters

5k+ Nodes

Query

10M / Day

Storage

Total: 100 PB  
Daily: 1PB

Peak IO

100M rows / s



腾讯广告



微信支付



腾讯视频  
V.QQ.COM



QQ

TENPAY.COM  
财付通



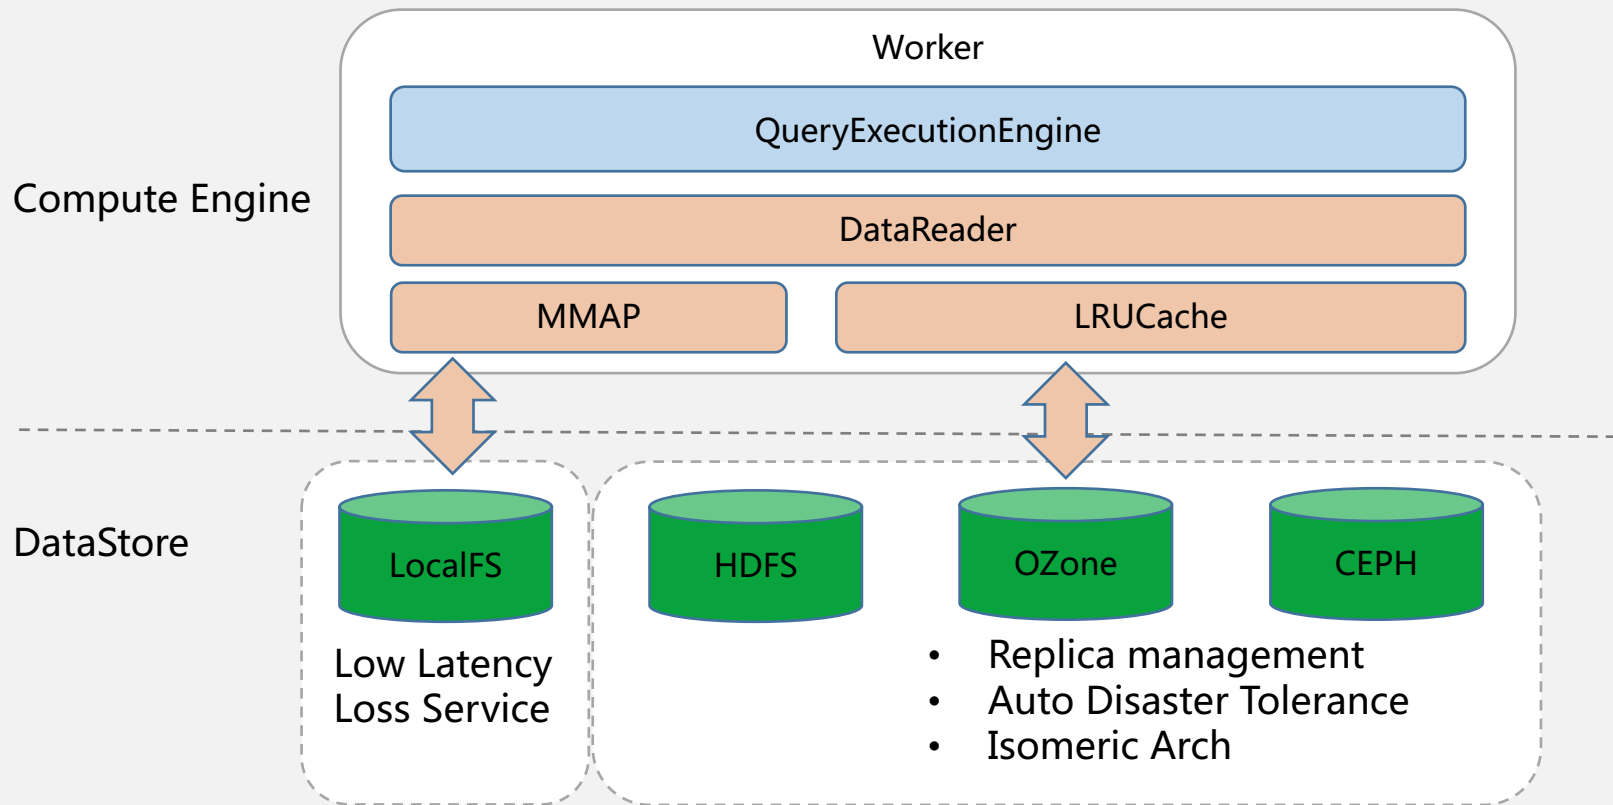
腾讯游戏  
Tencent Games

用心创造快乐!

# CONTENTS

- **01** | Background
- **02** | Storage: Various Columns and Indexes
- **03** | Computation: Integrated with Presto
- **04** | Benchmark and Applications

# Basic Arch



- Districted **vs** Local
- HA **vs** Low Latency
- Hot Data **vs** Cold Data

# Column-Oriented & Index

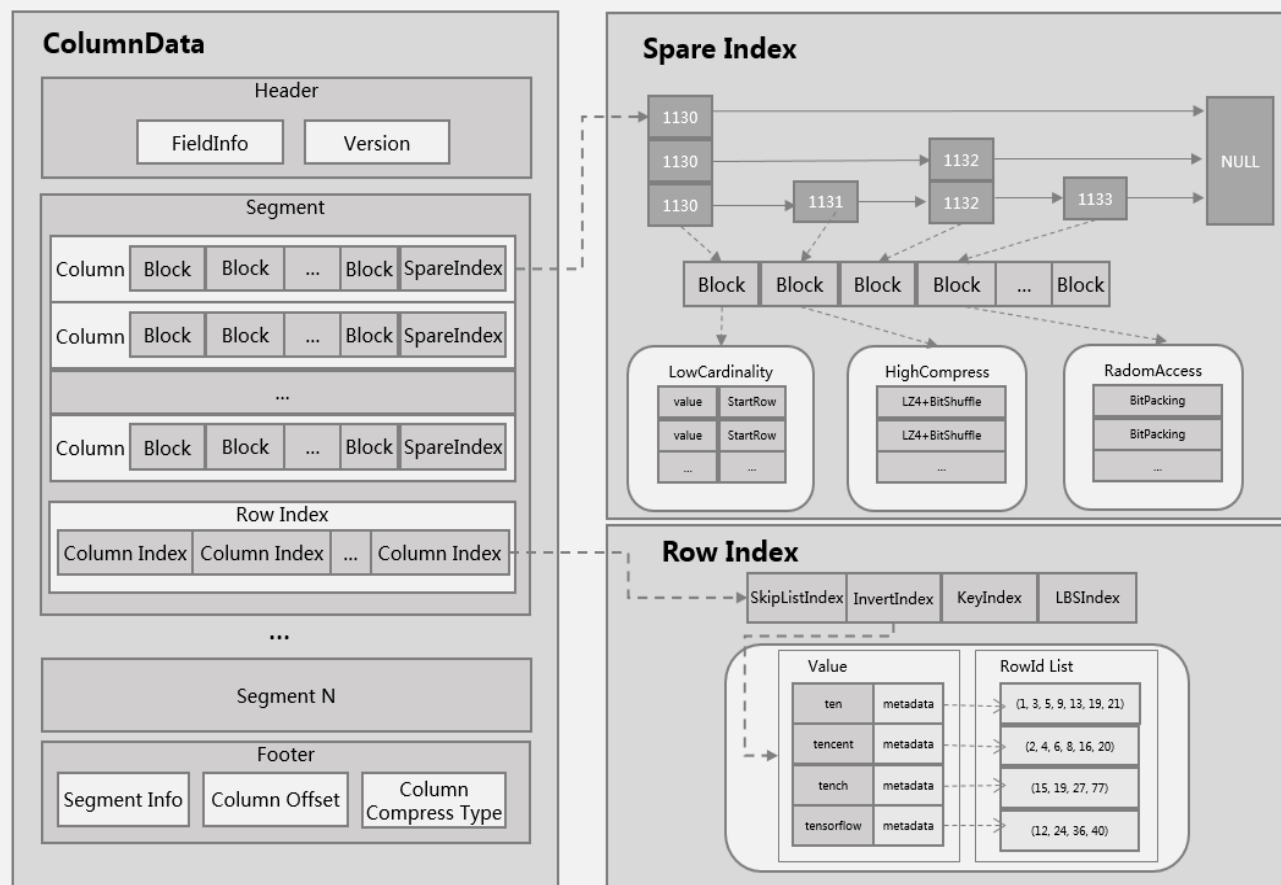
Q1: High QPS?  
Q2: Second response for 10Billion rows?  
Q3: Cost-friendly for mass data?

## Columns:

1. Retrieval —— Low latency
2. Sorted —— Mass Data
3. Compressed —— Cost-Friendly
4. Nested —— Support parquet

## Indexes:

1. SpareIndex
2. SkipListIndex
3. InvertedIndex
4. KeyIndex
5. LBSIndex





# Retrieval Column

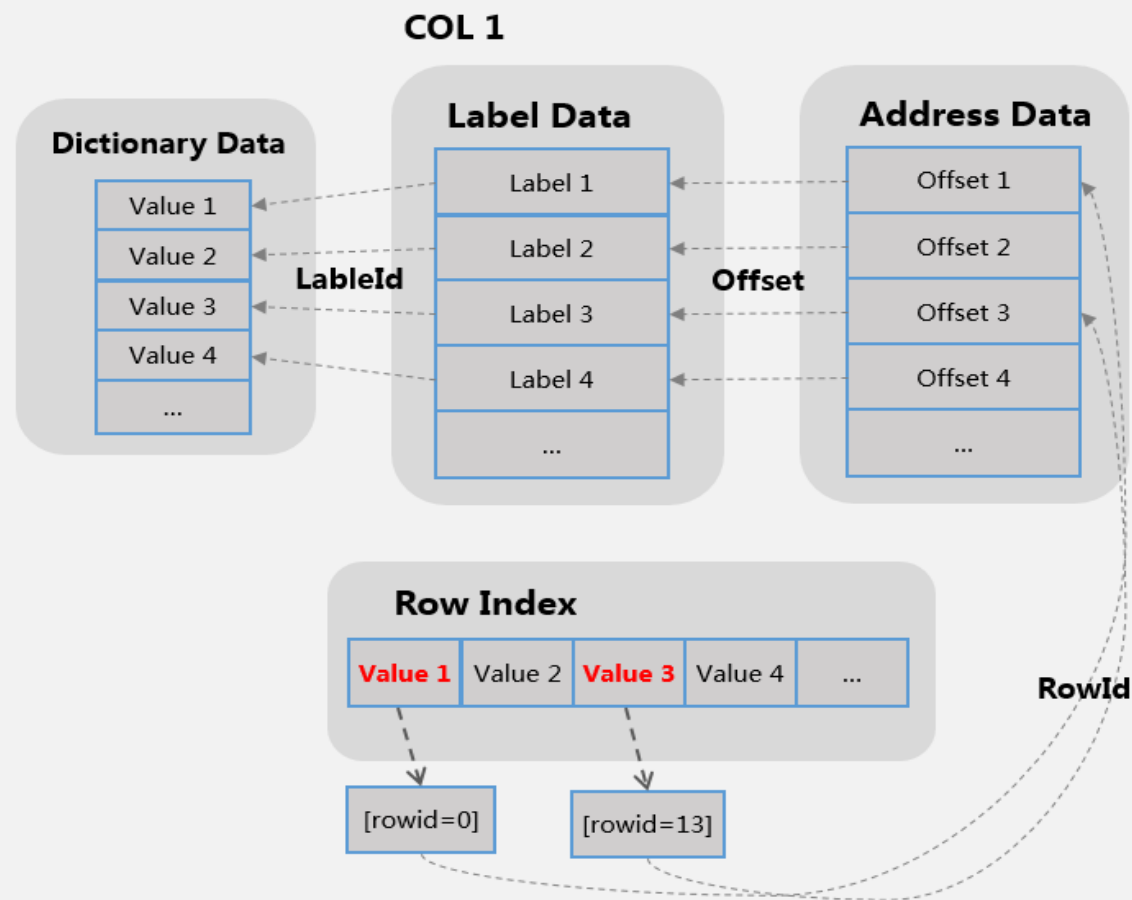
- **Application**

1. Low latency
2. Medium data
3. Simple Query

- **Implementation**

1. Storage -> Time
2. Dictionary Index

Index Size/Origin Data = 40%



# Sorted Column

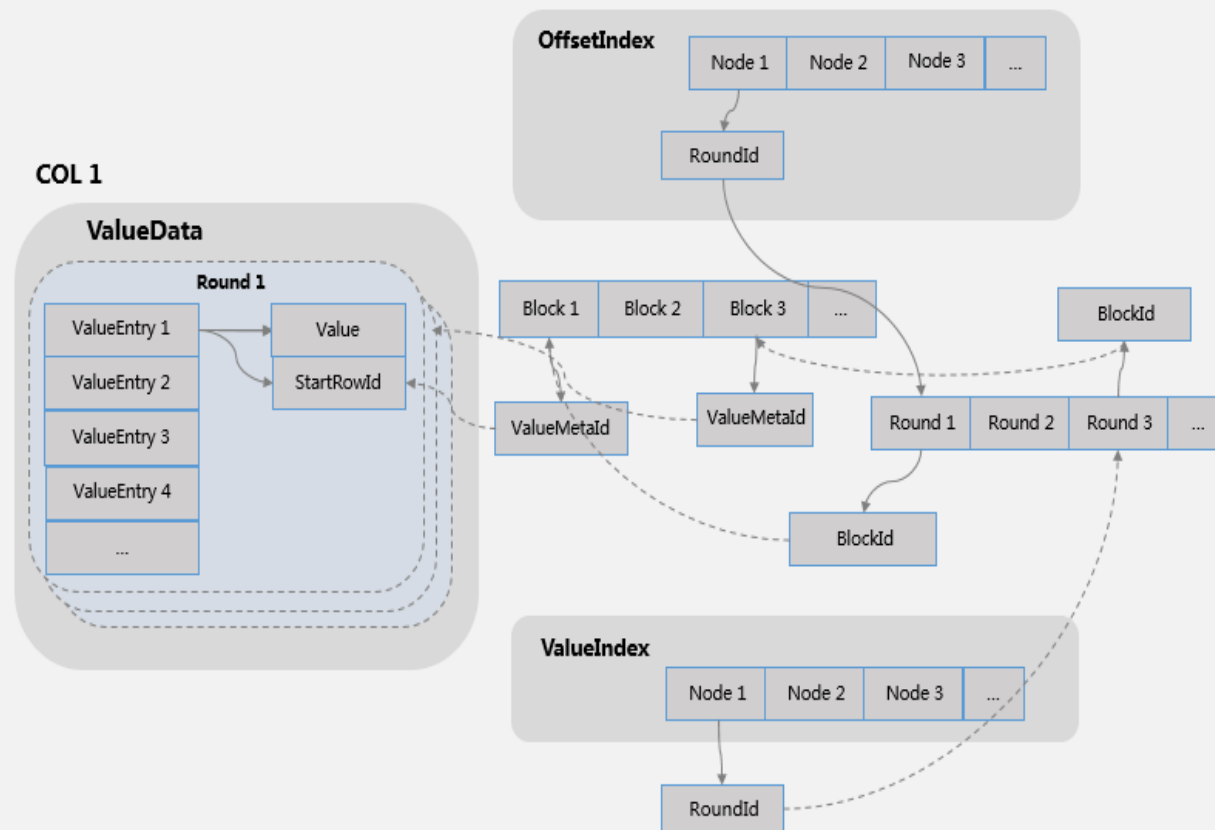
- **Application**

1. Mass Data (much bigger than memory)
2. No other accelerations

- **Implementation**

1. Sorted
2. Index both on data and offset

Improvement: 10x speed up



# Compressed Column



- **Application**

1. High Cardinal
2. Direct compression: low performance

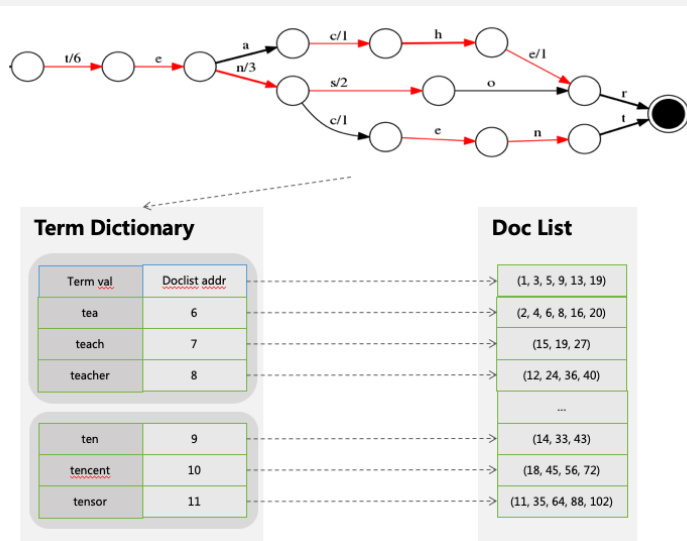
- **Implementation**

1. BitShuffle + LZ4

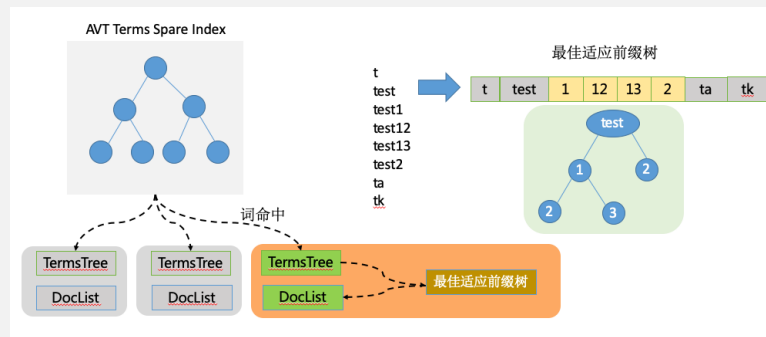
Compression rate: 50% up



## Inverted Index

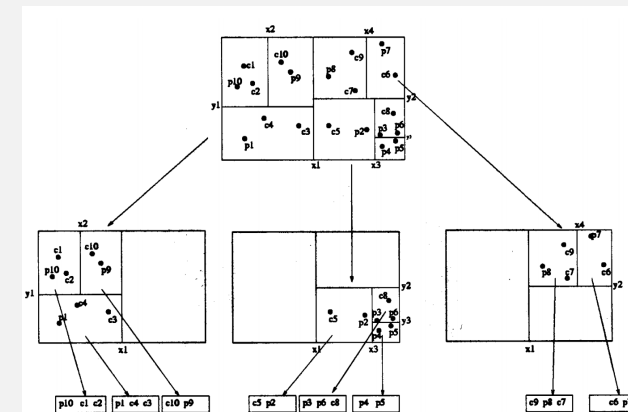


## FST KeyIndex



	完全基于 FST 的索引的 key 值索引	基于平衡二叉树的 key 值索引	Lucene 倒排索引
十万次数据点查询时	490 ms	243 ms	736 ms
相对于 Lucene 倒排索引的存储效果	1.3 倍	1.013 倍	基准 1
千次随机范围查询耗时	142706 ms	13737 ms	24163 ms

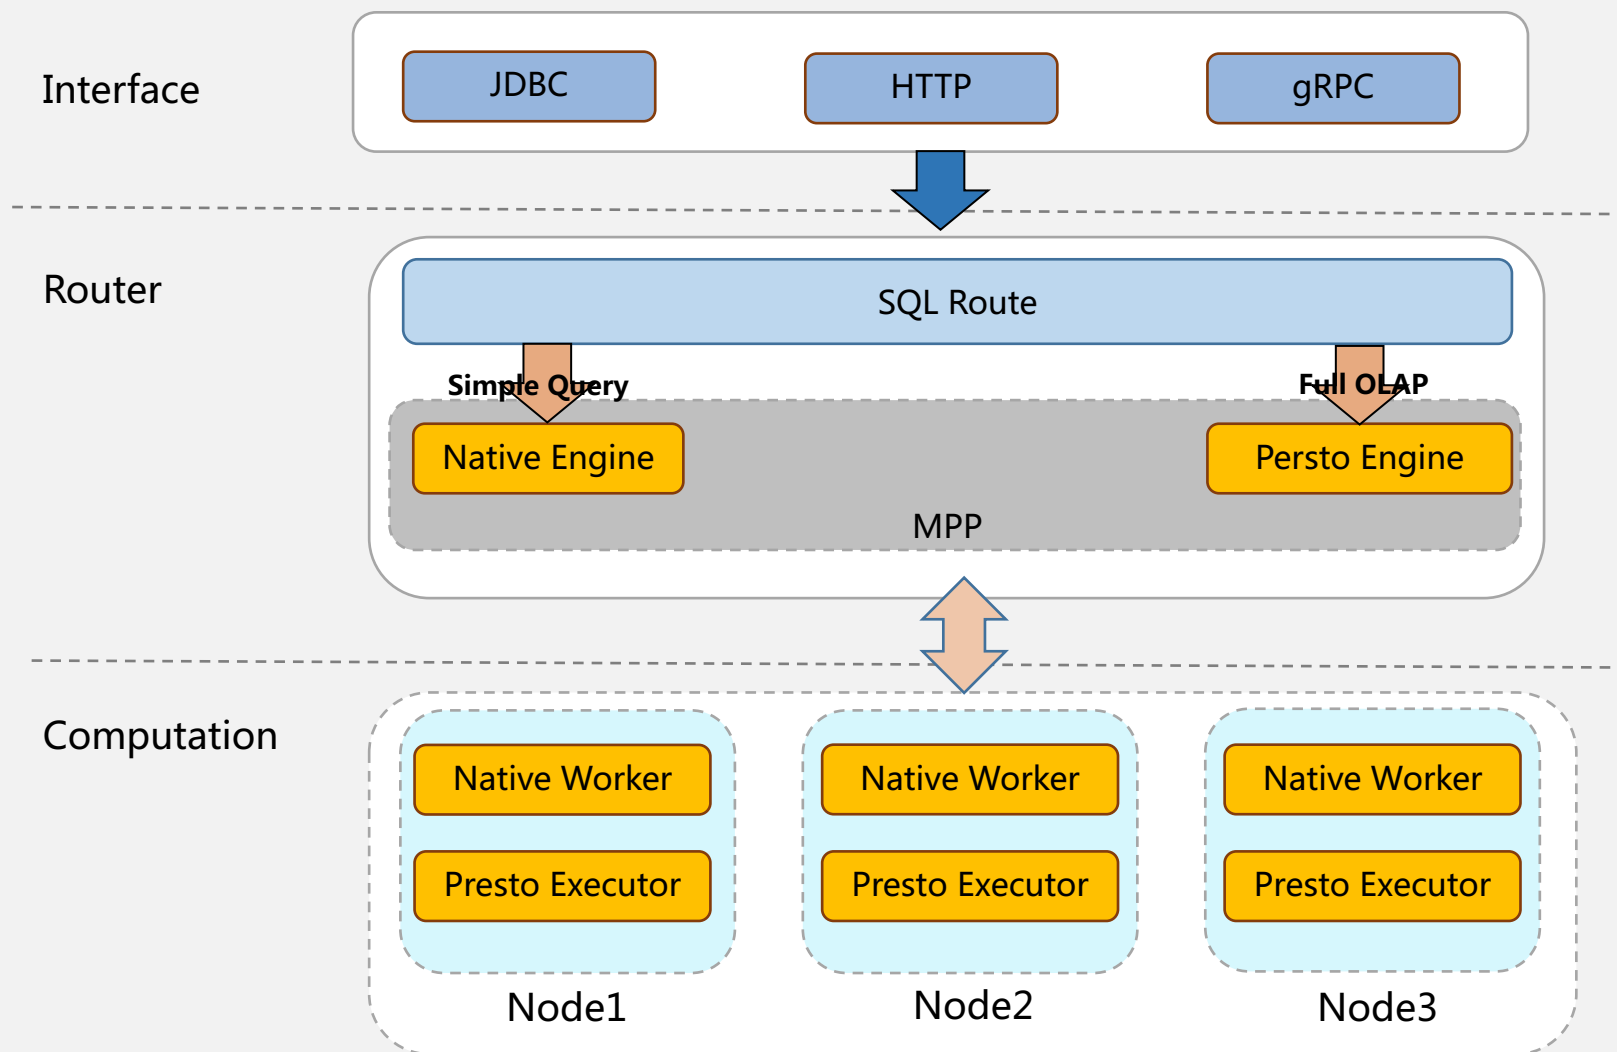
## LBS(KDB) Index



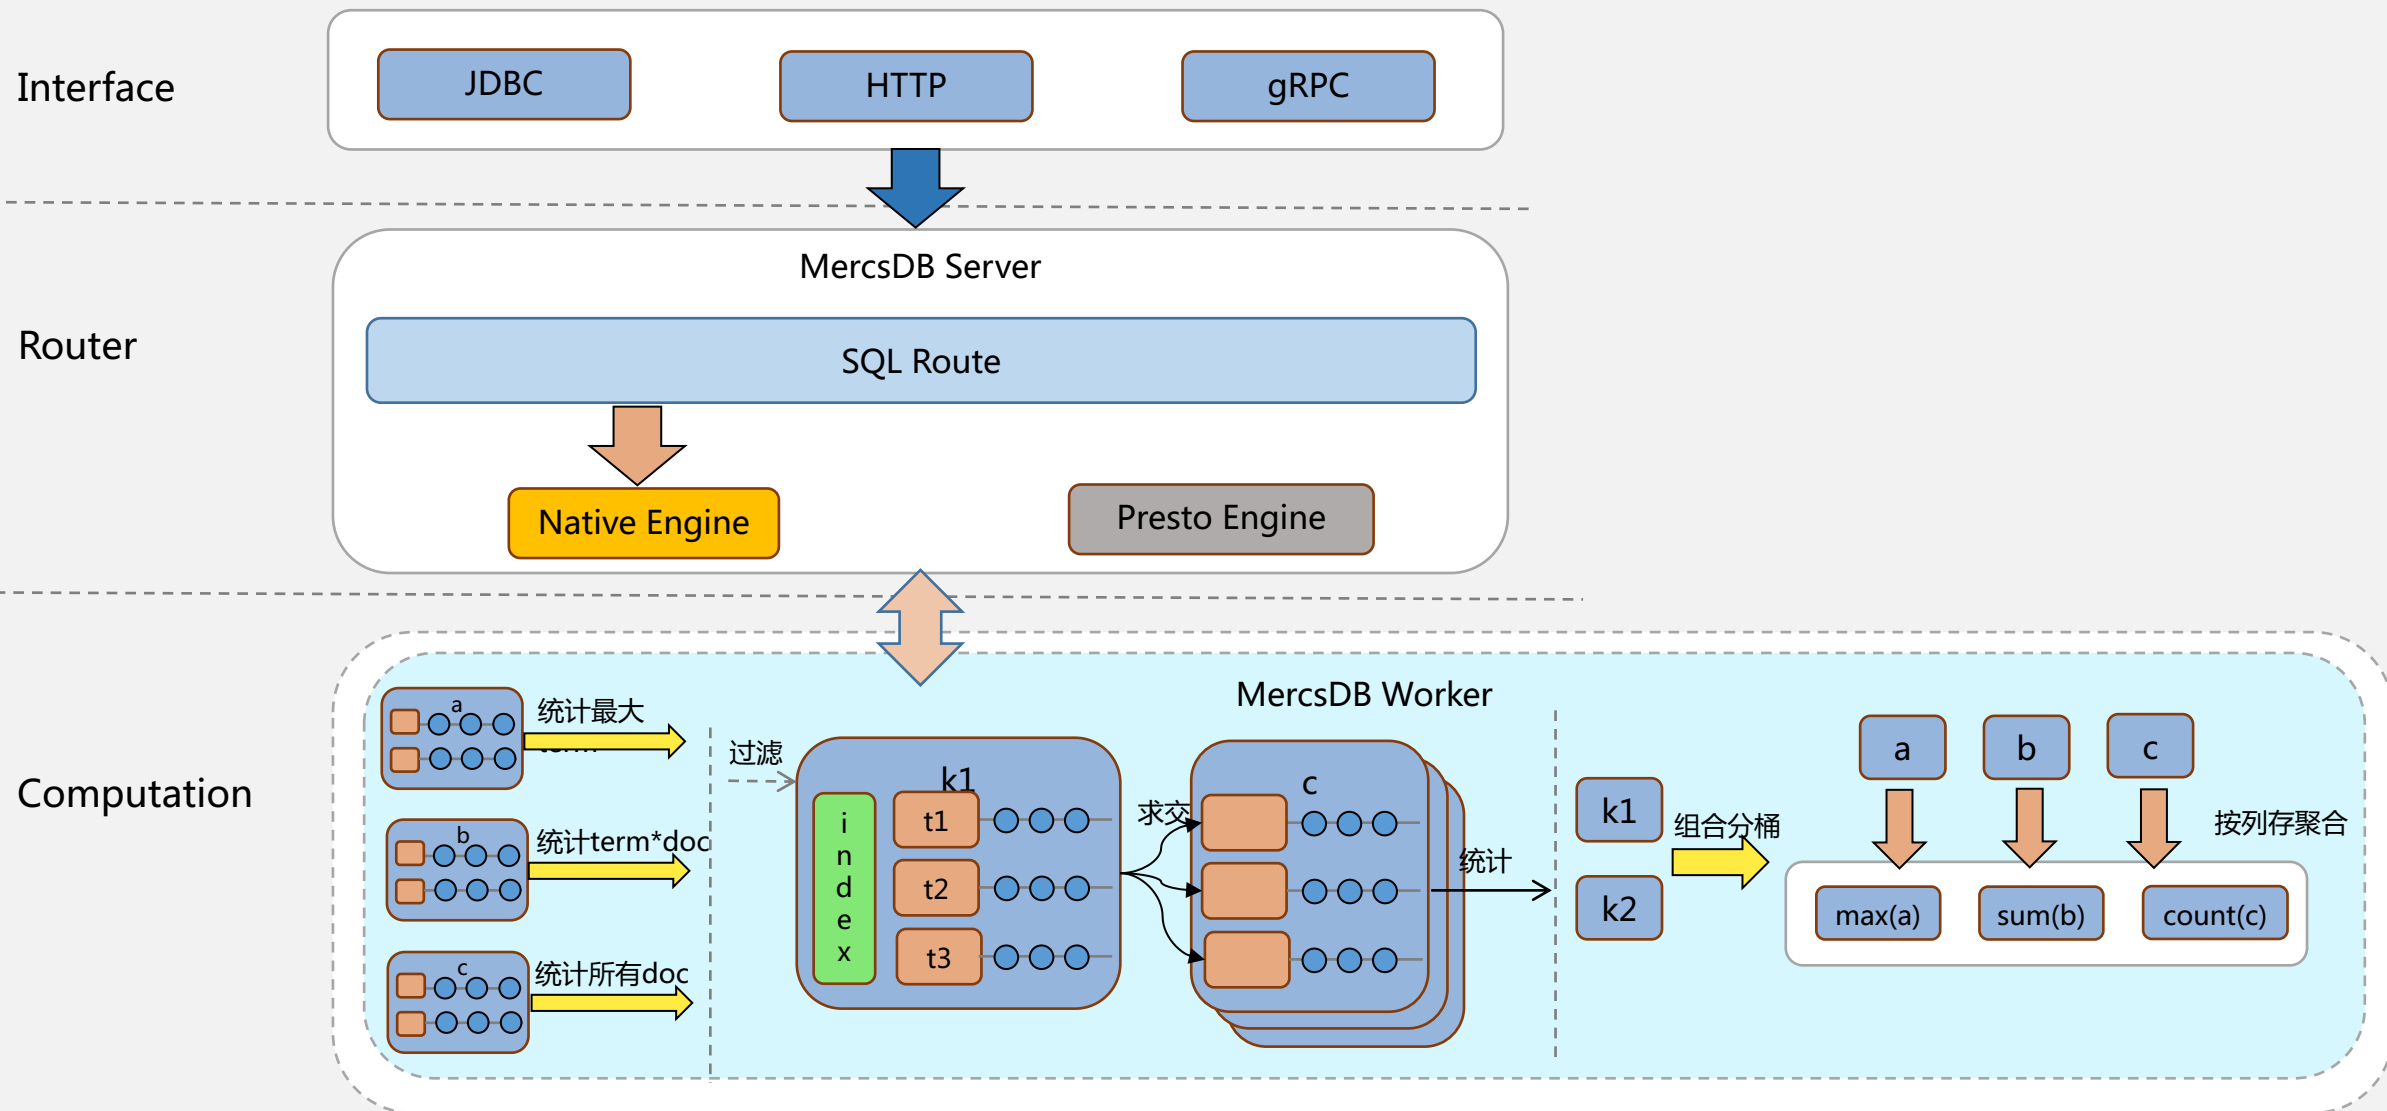
# CONTENTS

- **01** | Background
- **02** | Storage: Various Columns and Indexes
- **03** | Computation: Integrated with Presto
- **04** | Benchmark and Applications

# MercsDB — Dynamic Router



# MercsDB Native Engine

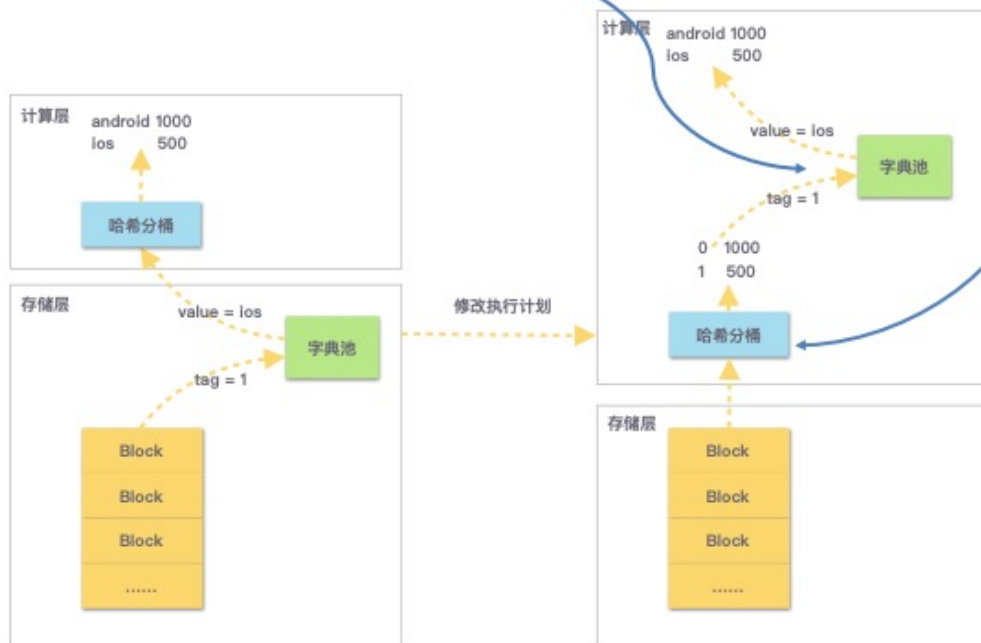


# Late Materialization

```
SELECT os, count(*) FROM table GROUP BY os
```

```
SELECT hermesTag(os_tag), count(*) FROM table GROUP BY hermesTag(os_tag)
```

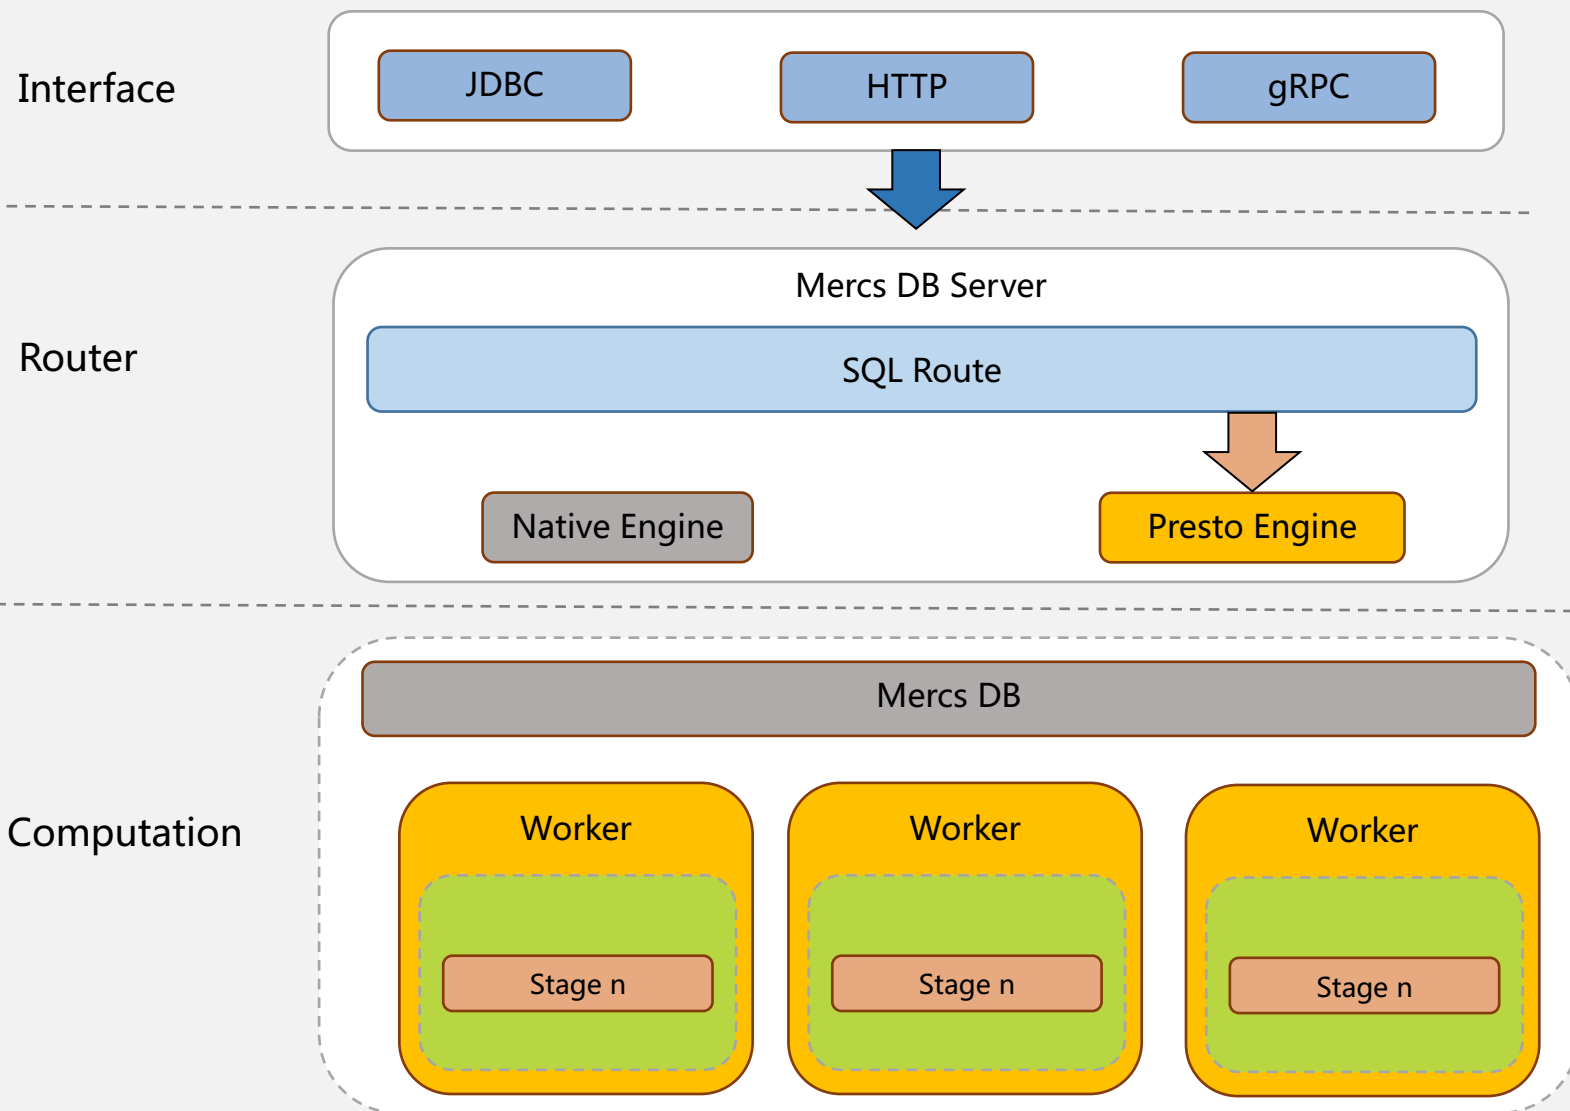
```
SELECT hermesTag(os_tag), c FROM (SELECT os_tag, count(*) as c FROM table GROUP BY os_tag)
```



便于理解按照SQL来表达转化,  
实际是执行计划中算子的替换  
大大提高了分桶计算的效率





# MericsDB Presto Engine

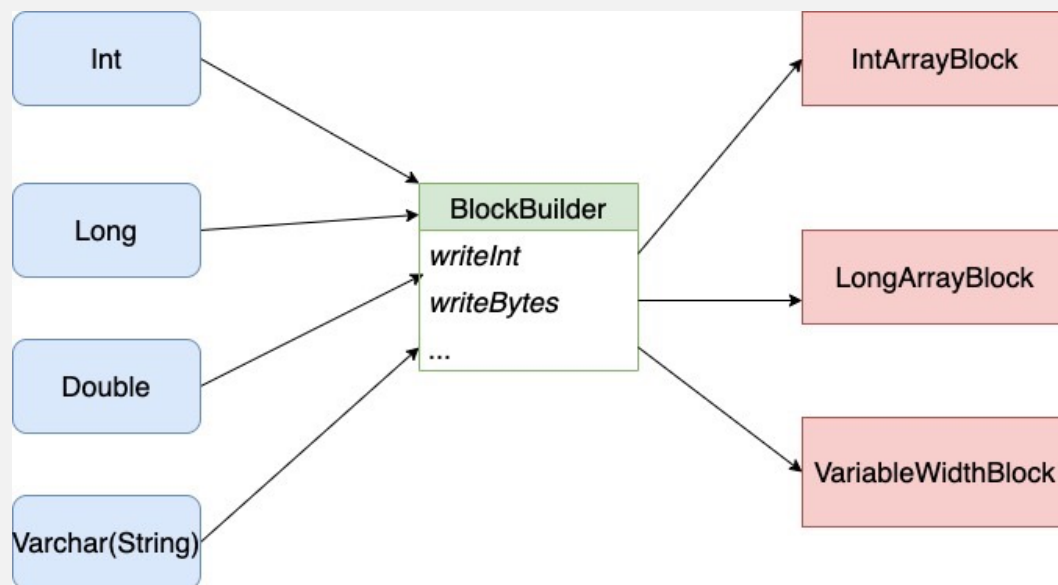


# Push down into MercsDB


SELECT \* FROM  
 (SELECT FROM A)  
 JOIN  
 (SELECT agg FROM B WHERE X)  
 ON C = D



# Data transformation



Presto Block	MercsDB Data
long[] Values	Result of Mercs API
bool[] isNulls	Vectorized

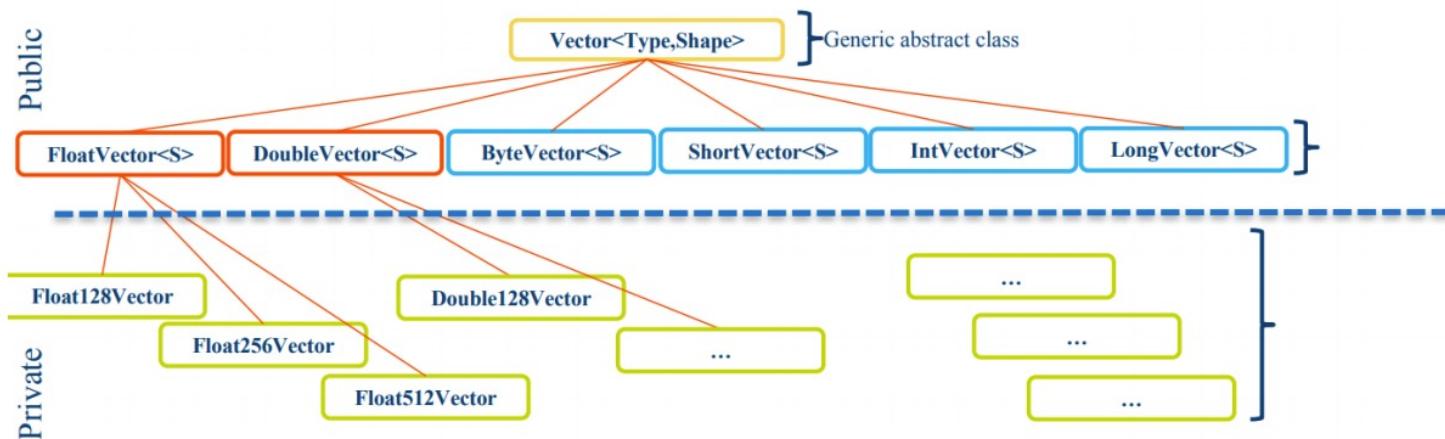


# Java Vector API

- Incubator since JDK 16
- Tencent [KONA JDK 17](#)

```
void scalarComputation(float[] a, float[] b, float[] c) {
    for (int i = 0; i < a.length; i++) {
        c[i] = (a[i] * a[i] + b[i] * b[i]) * -1.0f;
    }
}
```

## Vector API

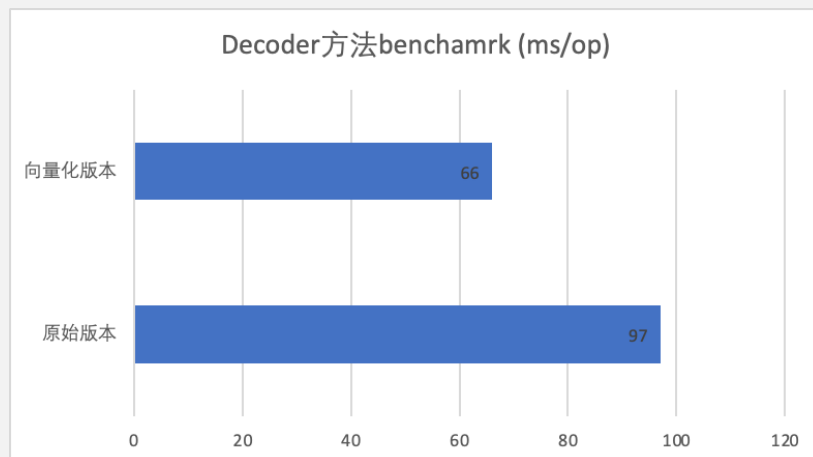


```
static final VectorSpecies<Float> SPECIES = FloatVector.SPECIES_PREFERRED;

void vectorComputation(float[] a, float[] b, float[] c) {
    int i = 0;
    int upperBound = SPECIES.loopBound(a.length);
    for (; i < upperBound; i += SPECIES.length()) {
        // FloatVector va, vb, vc;
        var va = FloatVector.fromArray(SPECIES, a, i);
        var vb = FloatVector.fromArray(SPECIES, b, i);
        var vc = va.mul(va)
                    .add(vb.mul(vb))
                    .neg();
        vc.intoArray(c, i);
    }
    for (; i < a.length; i++) {
        c[i] = (a[i] * a[i] + b[i] * b[i]) * -1.0f;
    }
}
```

# Vectorization with Vector API

- Opt1: Unroll Loop
- Opt2: Unify Vector Species
- Opt3:
  - No Boxing & Unboxing
  - No Object creation
  - No function call



```
void decode(long[] input, long[] output) {
    for (int i = 0; i < input.length; ++i) {
        long index = input[i];
        long offset = (index * 20) >>> 3;
        int v = readInt(offset) >>> 8;
        int shift = (int) ((index + 1) & 1) << 2;
        output[i] = (v >>> shift) & 0xFFFFF;
    }
}
```

```
static void decodeVector(long[] input, long[] output) {
    int i;
    int len = input.length - input.length % (4 * LONG_SPECIES.length());


    // main loop
    for (i = 0; i < len; i += 4 * LONG_SPECIES.length()) {
        LongVector v1 = LongVector.fromArray(LONG_SPECIES, input, i);
        v1.lanewise(VectorOperators.LSHL, 4).add(v1.lanewise(VectorOperators.LSHL, 2))
            .lanewise(VectorOperators.LSHR, 3).add(currentOffset).intoArray(values, offset: 0);
        LongVector v2 = LongVector.fromArray(LONG_SPECIES, input, offset: i + LONG_SPECIES.length());
        v2.lanewise(VectorOperators.LSHL, 4).add(v2.lanewise(VectorOperators.LSHL, 2))
            .lanewise(VectorOperators.LSHR, 3).add(currentOffset).intoArray(values, LONG_SPECIES.length());
        LongVector v3 = LongVector.fromArray(LONG_SPECIES, input, offset: i + 2 * LONG_SPECIES.length());
        v3.lanewise(VectorOperators.LSHL, 4).add(v3.lanewise(VectorOperators.LSHL, 2))
            .lanewise(VectorOperators.LSHR, 3).add(currentOffset).intoArray(values, offset: 2 * LONG_SPECIES.length());
        LongVector v4 = LongVector.fromArray(LONG_SPECIES, input, offset: i + 3 * LONG_SPECIES.length());
        v4.lanewise(VectorOperators.LSHL, 4).add(v4.lanewise(VectorOperators.LSHL, 2))
            .lanewise(VectorOperators.LSHR, 3).add(currentOffset).intoArray(values, offset: 3 * LONG_SPECIES.length());

        for (int j = 0; j < 4 * LONG_SPECIES.length(); j++) {
            values[j] = readInt(values[j]);
        }


        LongVector.fromArray(LONG_SPECIES, values, offset: 0)
            .lanewise(VectorOperators.LSHR, 8)
            .lanewise(VectorOperators.LSHR, v1.add(1).and(1).lanewise(VectorOperators.LSHL, 2))
            .and(0xFFFFF).intoArray(output, i);
        LongVector.fromArray(LONG_SPECIES, values, LONG_SPECIES.length())
            .lanewise(VectorOperators.LSHR, 8)
            .lanewise(VectorOperators.LSHR, v2.add(1).and(1).lanewise(VectorOperators.LSHL, 2))
            .and(0xFFFFF).intoArray(output, offset: i + LONG_SPECIES.length());
        LongVector.fromArray(LONG_SPECIES, values, offset: 2 * LONG_SPECIES.length())
            .lanewise(VectorOperators.LSHR, 8)
            .lanewise(VectorOperators.LSHR, v3.add(1).and(1).lanewise(VectorOperators.LSHL, 2))
            .and(0xFFFFF).intoArray(output, offset: i + 2 * LONG_SPECIES.length());
        LongVector.fromArray(LONG_SPECIES, values, offset: 3 * LONG_SPECIES.length())
            .lanewise(VectorOperators.LSHR, 8)
            .lanewise(VectorOperators.LSHR, v4.add(1).and(1).lanewise(VectorOperators.LSHL, 2))
            .and(0xFFFFF).intoArray(output, offset: i + 3 * LONG_SPECIES.length());
    }
    // ...
}
```

# Others

- Batch Vectorization



- Sequential memory access



# CONTENTS

- **01** | Background
- **02** | Storage: Various Columns and Indexes
- **03** | Computation: Integrated with Presto
- **04** | Benchmark and Applications


# SSB Benchmark 1

600M rows


Rows	Origin Size	MercsDB (with index)	ClickHouse (with spare and primary-key index)
600M	200GB	89GB	60GB
6B	2TB	882GB	593GB

60B rows

6亿数据查询耗时(ms)



60亿数据查询耗时(ms)



MercsDB




Optimized MercsDB



# SSB Benchmark 2: QPS 1

600M rows




CK




MercsDB

60B rows



# SSB Benchmark 3: QPS 20

600M rows




CK



MercsDB

60B rows



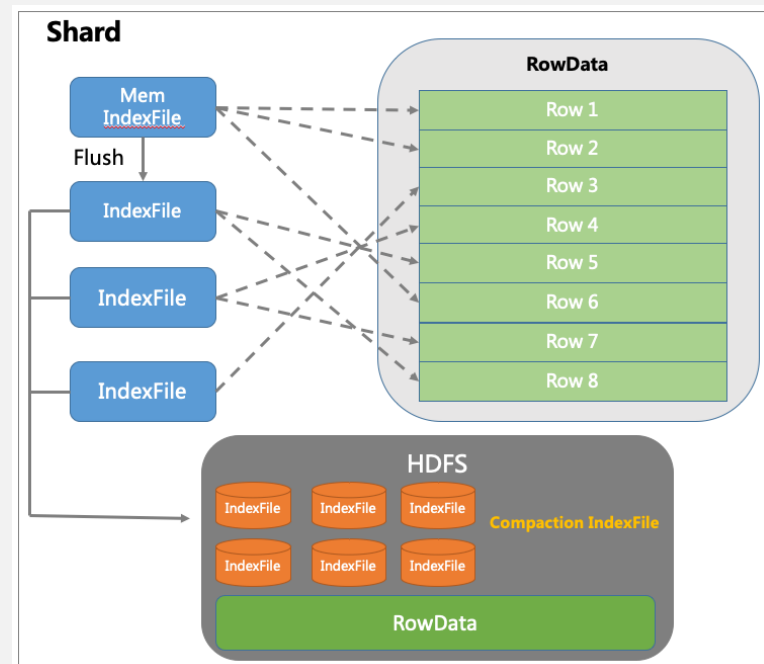
# Log Retrieval in WeChat Pay

## ●Background:

1. Real Time Write: 100 Billion Rows / day
2. Retrieval both Global and Specific
3. Mass Data

## ●Solutions:

1. Write with TubeMQ
2. Participle and Index on Data
3. Separation of storage:
  - I. IndexFile: Local Disk
  - II. RowData: HDFS



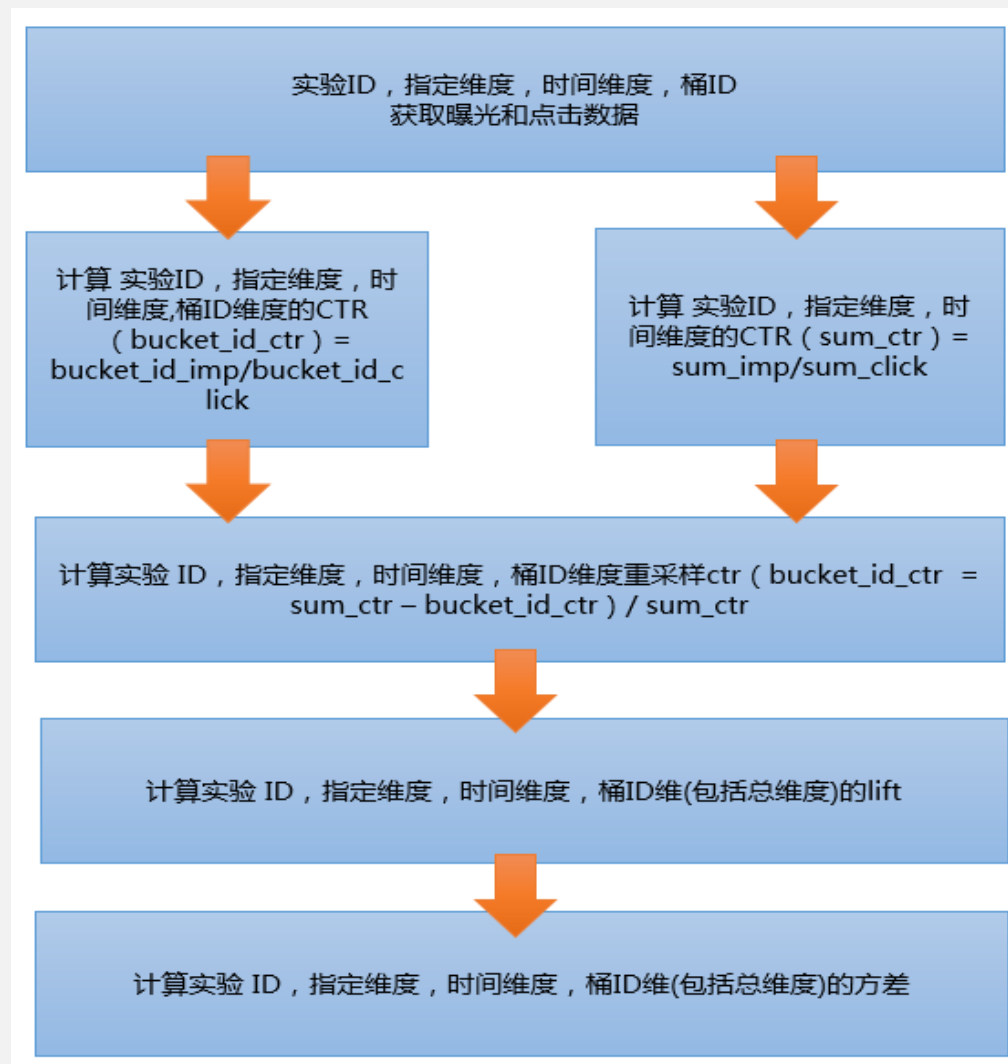
# AB Test for Ads

## ●Background:

1. Second response for Mass Data
2. Thousand columns in single query
3. Arbitrary Join


## ●Solutions:

1. Sorted by Primary Key (ad\_id)
2. Compressions on metric column, 40% storage of origin data
3. Presto/Spark supported with Cache



# Future Plan

- Cloud Service & Open-Source
- Vectorization
- Fault-Tolerance
- Memory Management
- ...



腾讯 MercsDB 技术探讨



该二维码 7 天内 (6月17日前) 有效, 重新进入将更新

腾讯大数据



# Thanks



数据平台部